

# Weak References for EcmaScript

---

E. DEAN TRIBBLE – AGORIC SYSTEMS

MARK S. MILLER – GOOGLE, INC.



# Intended Audience

---

- The garbage collection challenges addressed here largely arise in the implementation of libraries and frameworks.
- The features proposed here are advanced features that are primarily intended for use by library and framework creators, not their clients.
- Thus, the priority is enabling library implementors to correctly, efficiently, and securely manage object lifetimes and finalization.

# Weak References and Finalization

---

- Fundamental expressiveness
- needed for libraries and frameworks
- to simplify memory management
- and support...
  - Remote references
  - Observers for MVC and data-binding
  - DOM iterators
  - Reactive-style libraries
  - Handles for wasm/external resources

Smalltalk

Java

C#

E

Lua

Haskell

Python

Racket

...

# Requirements for Approach

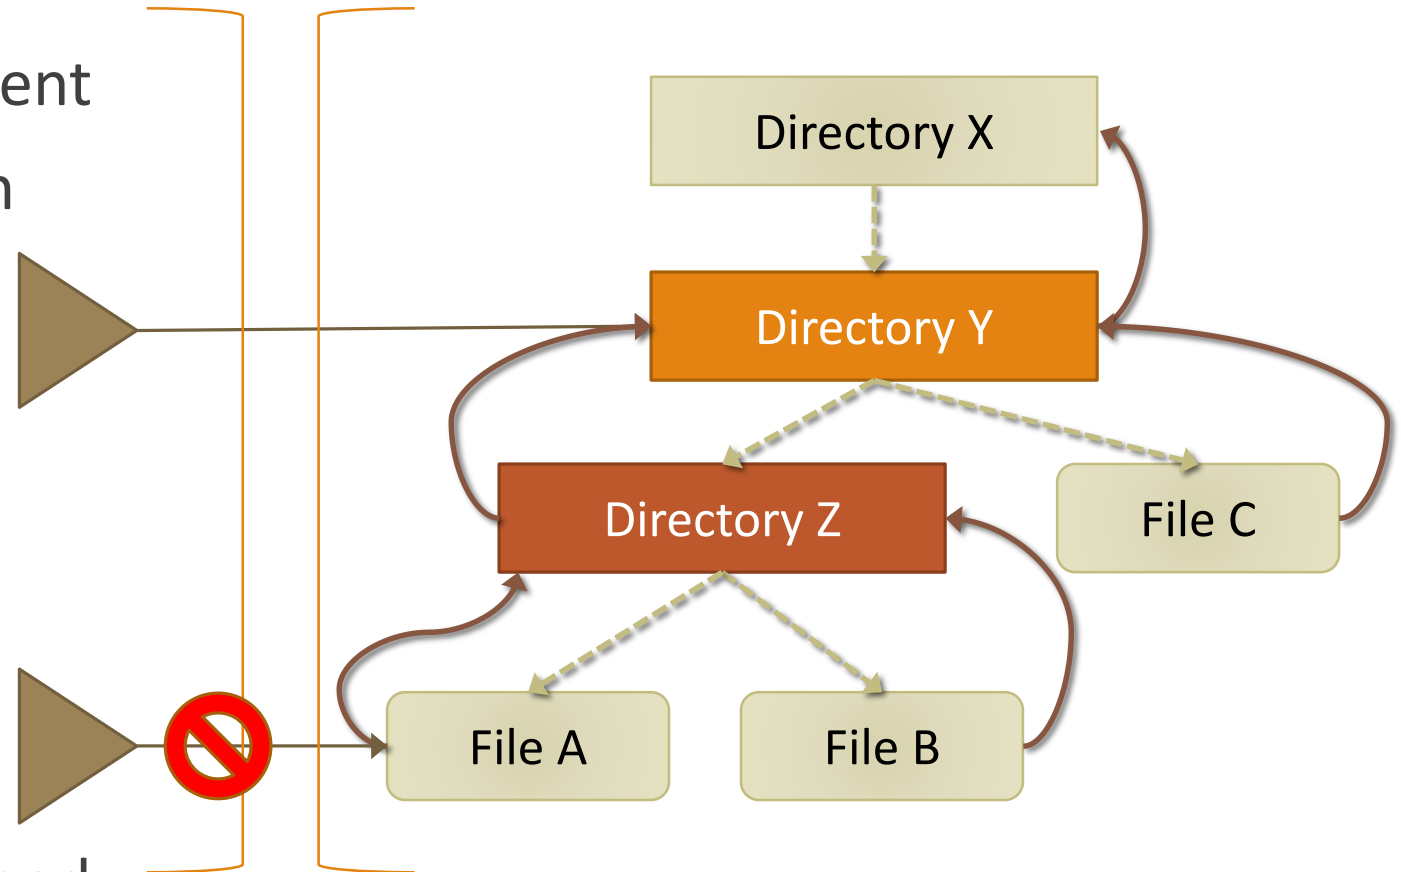
---

- Multiple, internal and external finalization
  - Remote reference – do something when **I** go away
  - Observers – do something when **you** go away
- Preclude Resurrection
  - Condemned object becomes reachable again
  - But its world is broken
- Avoid Layered collection
  - Data structure is reclaimed one layer at a time

# Resurrection Damage

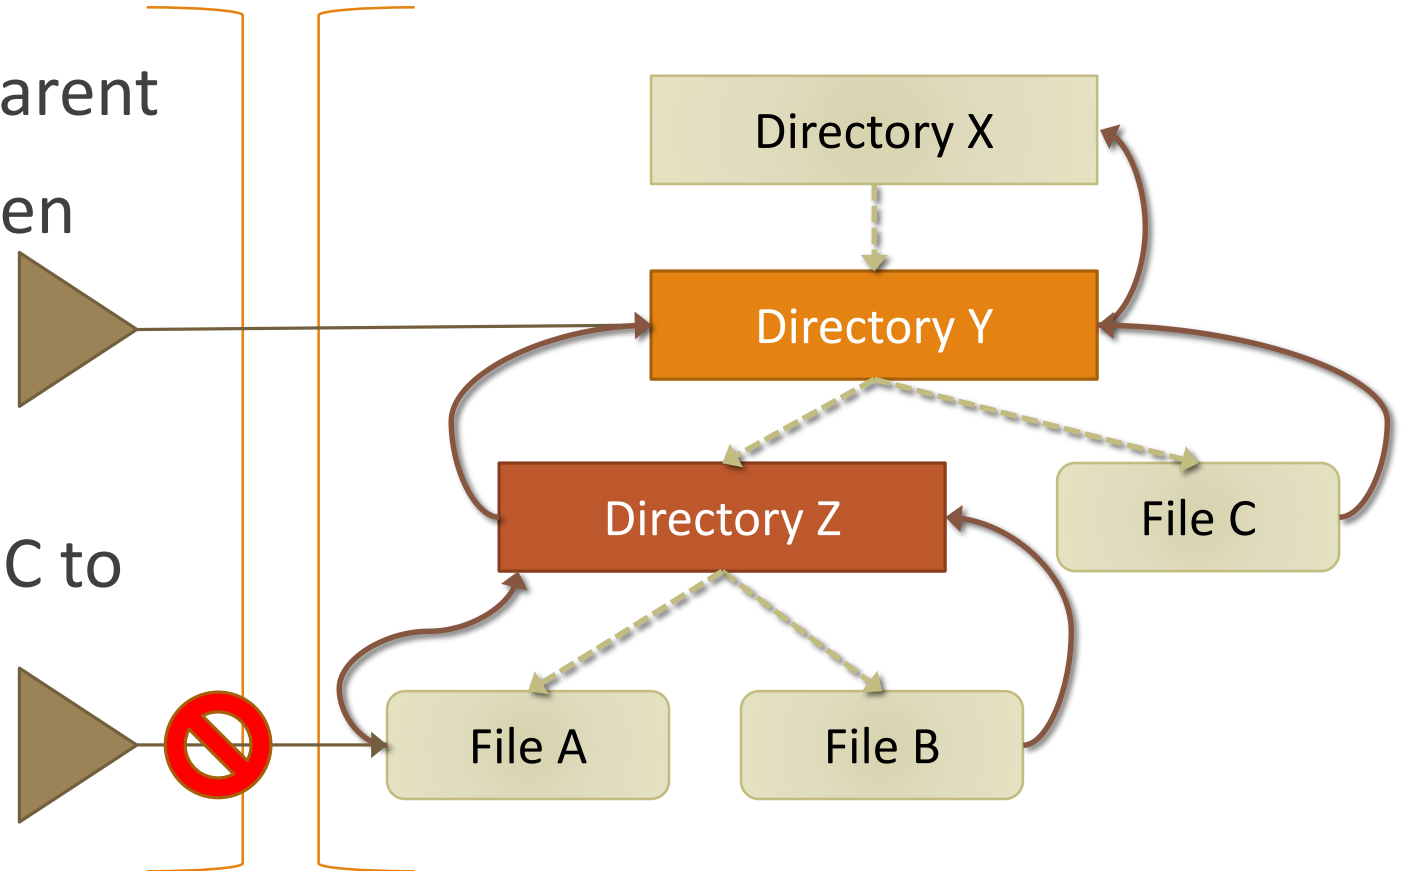
- Files/directories point at parent
- Parent canonicalizes children

- A gets finalized first
- Z is made reachable again
- HAZARD: The Z tree is damaged



# Layered Collection

- Files/directories point at parent
- Parent canonicalizes children
- Cleanup at each layer
- Don't require a separate GC to reclaim each layer



# Terminology

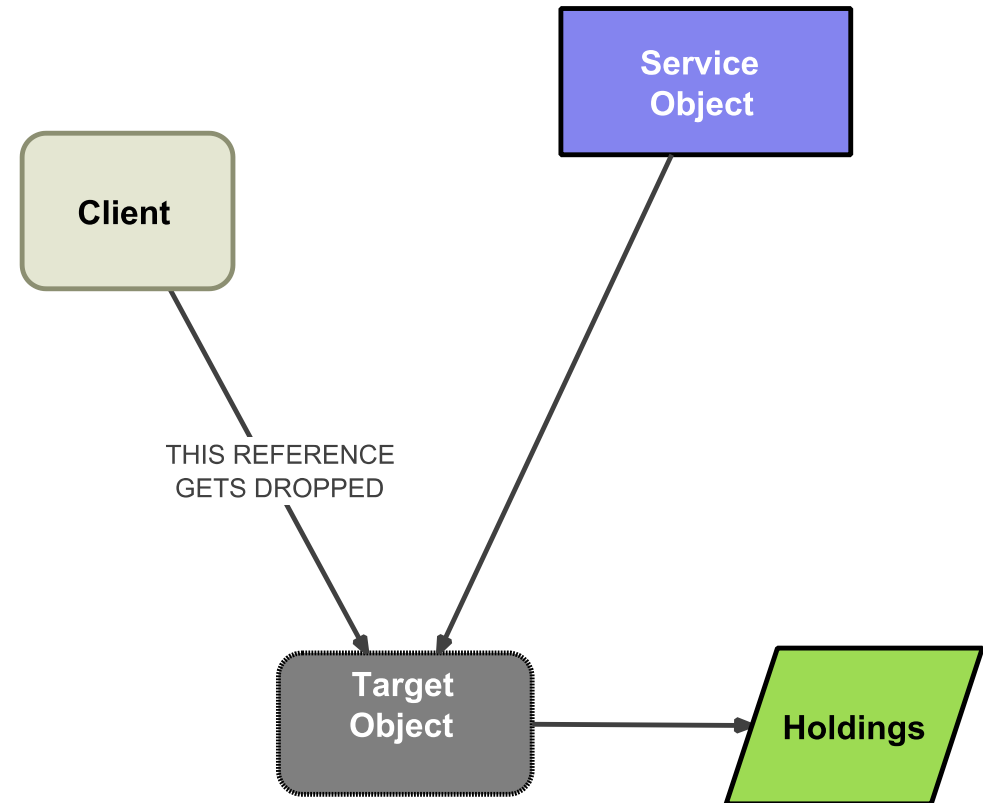
---

- **Reachable** objects – objects that program execution can reach
- **Strongly Reachable** objects – objects that program execution can reach without dereferencing weak references
- **Reclaimed** object – an object that the garbage has noticed is not strongly reachable and has made not reachable
- **Weak reference** – allows access to an object that has not yet been reclaimed, but does not prevent that object from being reclaimed
- **Finalization** – the execution of code to clean up after an object that has been reclaimed
- **Conservative** – some objects that are not strongly reachable may never be reclaimed

# Basic Scenario

---

- Client gets target from the service
- Target uses holdings to accomplish it's function
- When client drops target, service wants to clean it up
- Service retains Target indefinitely!

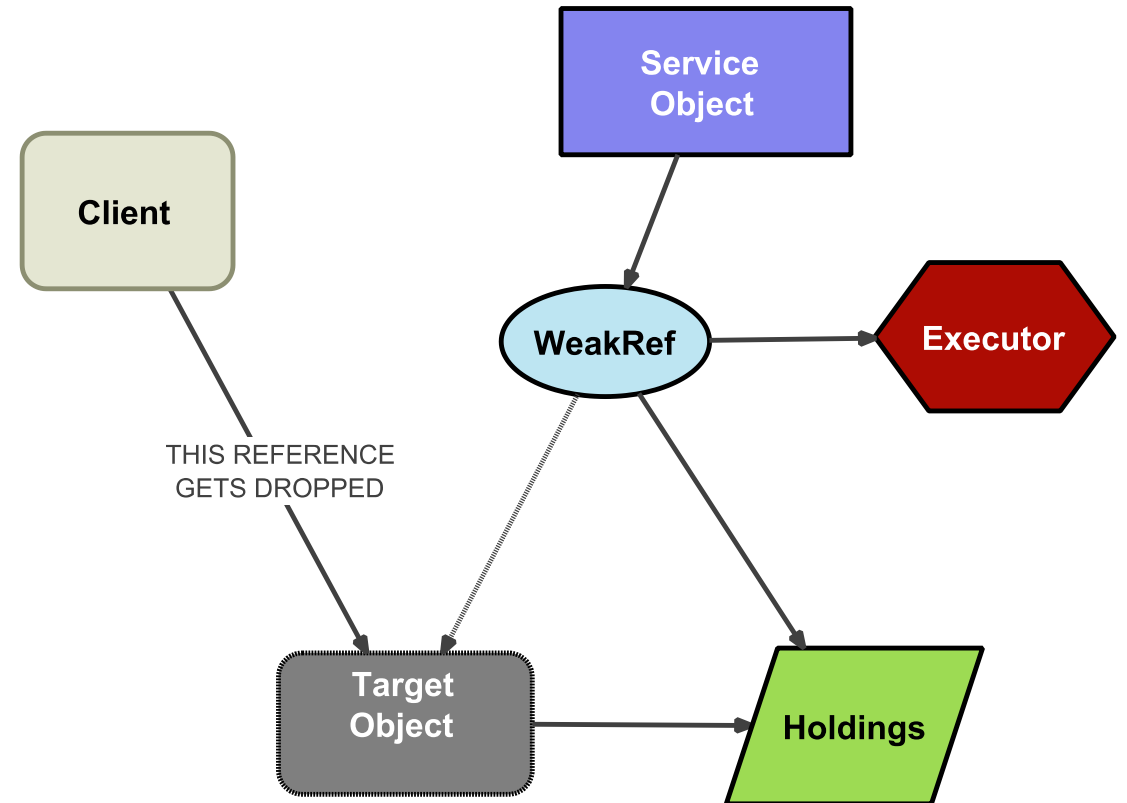




# Basic Approach

---

- Insert a WeakRef
- Weak reference shares holdings
- Service's cleanup is scheduled when Target is reclaimed
- The cleanup uses Holdings



# Proposed API

---

- **WeakRef**
  - Points weakly at a Target
  - Holds an associated value that gets used during cleanup for a Target
- **WeakRefGroup**
  - Create `WeakRefs` for a related group of Targets
  - Manage cleanup for `WeakRefs` whose Targets have been reclaimed

# class **WeakRef** {

---

- **deref ( )**
  - Return a strong reference to the weakly-held Target object, or undefined if it has been reclaimed
  - The Target will be retained until the end of the turn
- **drop ( )**
  - Drop the weak reference to the Target and prevent finalization
- **get holdings ( )**
  - Return the holdings associated with the Target

# class WeakRefGroup {

---

- **constructor**(cleanup = undefined) => WeakRefGroup
  - cleanup(weakRefs) => void
    - Invoked in a new job to cleanup after reclaimed targets
    - weakRefs – iterator of weakRefs whose Targets are reclaimed
- **makeRef**(target, holdings = undefined) => WeakRef
  - target – object pointed to weakly; returned by deref()
  - holdings – target-specific arg used for cleanup after the Target (optional)
- **shutdown**() => void
  - The subsystem is being shutdown; perform no cleanup action for any WeakRef in this group
- **purgeSome**(cleanupNow) => void

# WeakRef States

---

- Available
  - Target is accessible
- Dirty
  - Target is inaccessible but cleanup is needed
- Clean
  - Target is inaccessible and cleanup is NOT needed

# Example: Observer without finalization

---

- Model points weakly to observers
- Doesn't really cleanup after GC'd observers

```
class Model {  
    #observers, #wrGroup;  
    constructor() {  
        this.#observers = new Set();  
        this.#wrGroup = new WeakRefGroup();  
    }  
  
    addObserver(observer) {  
        const wr = this.#wrGroup.makeRef(observer);  
        this.#observers.add(wr);  
    }  
    notify(msg) {  
        this.#observers.forEach(wr =>  
            wr.deref() && wr.deref().changed(msg));  
    }  
}
```

# Example: Observer using Finalization

---

- GC'd observers are removed
- No cleanup needed if the whole model is GC'd

```
class Model {  
  #observers, #wrGroup;  
  constructor() {  
    this.#observers = new Set();  
    const cleanup = iter => {  
      for (const wr of iter) {  
        this.#observers.delete(wr);  
      }  
    };  
    this.#wrGroup = new WeakRefGroup(cleanup);  
  }  
  addObserver(observer) {  
    const wr = this.#wrGroup.makeRef(observer);  
    this.#observers.add(wr);  
  }  
  notify(msg) {  
    this.#observers.forEach(wr =>  
      wr.deref() && wr.deref().changed(msg));  
  }  
}
```

# Example: Remote Reference

---

- Client-side reference to an object on the server
- Over a shared connection
- Send “DROP” when the client reference is GC’d
- No cleanup if the whole connection is GC’d

```
class RemoteConnection {
  #transport, #remotes, #wrGroup;
  constructor(transport) {
    this.#transport = transport;
    this.#remotes = new Map();
    const cleanup = iter => {
      for (const wr of iter) {
        this.dropRef(wr.holdings);
      }
    };
    this.#wrGroup = new WeakRefGroup(cleanup);
  }
  makeRemoteRef(remoteId) {
    const remoteRef = ...; // remoteRef construction elided
    const wr = this.#wrGroup.makeRef(remoteRef, remoteId);
    this.#remotes.set(remoteId, wr);
    return remoteRef;
  }
  dropRef(remoteId) {
    this.#transport.send("DROP", remoteId);
    this.#remotes.delete(remoteId);
  }
  ...
}
```



# Example: Connecting JS to wasm

---

- JS wrapper of an unmanaged wasm obj
- Bookkeeping by WasmConnection
- `wasmBridge.delete` when wrapper is GC'd
- No cleanup needed if wasm instance is GC'd

```
class WasmConnection {
  #wasmBridge, #wrappers, #cleanup, #wrGroup;
  constructor(wasmBridge) {
    this.#wasmBridge = wasmBridge;
    this.#wrappers = new Map();
    this.#cleanup = iter => {
      for (const wr of iter) {
        this.dropWrapper(wr.holdings);
      }
    };
    this.#wrGroup = new WeakRefGroup(cleanup);
  }
  makeWrapper(wasmAddr) {
    const wrapper = this.#wasmBridge.makeWrapper(wasmAddr);
    const wr = this.#wrGroup.makeRef(wrapper, wasmAddr);
    this.#wrappers.set(wasmAddr, wr);
    return wrapper;
  }
  dropWrapper(wasmAddr) {
    this.#wasmBridge.delete(wasmAddr);
    this.#wrappers.delete(wasmAddr);
  }
  ...
}
```

# class **WeakRefGroup** {

---

- **constructor**(cleanup = undefined) => WeakRefGroup
  - cleanupNow(weakRefs) => void
- **makeRef**(target, holdings = undefined) => WeakRef
- **shutdown**() => void
- **purgeSome**(cleanupNow) => void
  - cleanupNow(weakRefs) => void
    - Invoked synchronously to cleanup after some reclaimed targets
    - weakRefs – iterator of weakRefs whose Targets are reclaimed

# Example: JS ↔ long lived wasm

---

- WasmConnection encapsulates wrGroup and all its weakrefs.
- Only getWrapper observes wrapper availability.
- Long lived uni-directional uses purge, but never observe.

```
class WasmConnection {
  #wasmBridge, #wrappers, #cleanup, #wrGroup;
  constructor(wasmBridge) { ... }
  makeWrapper(wasmAddr) { ... }
  dropWrapper(wasmAddr) { ... }

  // Only for bi-directional wrapper identity
  getWrapper(wasmAddr) {
    const wr = this.#wrappers.get(wasmAddr);
    return wr && wr.deref(); // Only call to deref
  }

  // Only to cleanup early, before turn is over
  deleteSome() {
    this.#wrGroup.purgeSome(this.#cleanup);
  }
}
```

# Proposal Characteristics

---

- Automatic cleanup actions are scheduled in their own jobs
- Multiple, independent WeakRefs could have the same target
  - Internal finalization – remote reference
  - External finalization - observer
- Clients may unregister from finalization (`weakRef.drop()`)
- Read consistency

# Read Consistency

---

- The multiple-use hazard

```
weak.deref() && weak.deref().changed(msg);
```

```
if (observers.get(myKey)) {  
    ...do some expensive setup...  
    doOperation(myName, observers.get(myKey));  
}
```

- *A program cannot observe a target get reclaimed within the execution of a job (turn of the event loop)*
- Trivially precludes the multiple-use hazard
- Minimizes visible non-determinism

# Semantic Details

---

- WeakRef
  - Points weakly at Target
  - Points strongly at holdings
  - Points strongly at WeakRefGroup
- WeakRefGroup
  - Points strongly at cleanup function
  - Points strongly at dirty WeakRefs in group
- Cleanup iteration
  - Iterator is only productive during the call to cleanup
  - WeakRef is marked clean when it is pulled from the iterator

# Non-determinism in GC

---

- Weak references make GC behavior visible
- Racy-reads/writes present the same problems
  - It can appear sequentially consistent
  - If a program counts on that, it appears correct until it fails in production
- The non-determinism “bandwidth” is limited
- We can further mitigate it

# Minimize and contain non-determinism

---

- Deterministic computation advantages
  - Testability of ES components, frameworks, and applications
  - Reproducibility of results and bug reports
  - Portability across different runtimes and environments
  - Restricts who can read side-channels
- Read consistency
- Weak Reference construction should be closely held
  - `WeakRefGroup` is in the `System` namespace
- Cross-realm references are strong



# Open questions

---

- Are WeakRefs retained by their WeakRefGroup?
  - Fallback: Yes
- Is there a query operation for the states of a WeakRef
  - Fallback: No

Thank you

---

# WeakValueMap – with Leak

---

- Keys and Values are weak
- Entry remains after value is GC'd
  - LEAK!

```
class WeakValueMap {  
    constructor() {  
        this.map = new WeakMap();  
    }  
  
    get(key) {  
        let weakRef = this.map[key];  
        return weakRef && weakRef.get();  
    }  
  
    set(key, value) {  
        map.set(key, makeWeakRef(value));  
    }  
}
```

# WeakValueMap

---

- Keys and Values are weak
- Cleanup keys when value is GC'd
- Avoids storage leak

```
class WeakValueMap {  
    constructor() {  
        this.map = new WeakMap();  
        this.executor =  
            keyRef => this.map.delete(keyRef.get());  
    }  
  
    get(key) {  
        let weakRef = this.map[key];  
        return weakRef && weakRef.get();  
    }  
  
    set(key, value) {  
        let keyRef = makeWeakRef(key);  
        let valRef = makeWeakRef(v, this.executor, keyRef);  
        map.set(key, valRef);  
    }  
}
```

# More Terminology

---

- ***Condemned*** object – a (strongly) unreachable object that the garbage collector has noticed it can reclaim

# Read Consistency

---

- A target could be pointed to by multiple WeakRefs

# Unintended retention

---

```
const openfiles = new Map(); // file => weakRef(filestream)

const cleanupFile(file) {
  file.close();
  openFiles.delete(file);
  console.info("Filestream gc before close: ", file.name)
}

class FileStream {
  constructor(filename) {
    this.file = new File(filename, "r");
    openFiles.set(file, makeWeakRef(this, () => closeFile(this.file)));

    // now eagerly load the contents
    this.loading = file.readAsync().then(data => this.setData(data));
  } ...
}
```

- HAZARD: The executor function retains `this`

# Unintended retention – runtime black magic

---

```
class FileStream {  
  constructor(filename) {  
    let file = new File(filename, "r");  
    this.file = file;  
    openFiles.set(file, new WeakRef(this, () => closeFile(file)));  
  
    // now eagerly load the contents  
    this.loading = file.readAsync().then(data => this.setData(data));  
  } ...  
}
```

- The *unrelated* second function closes over `this`
- The runtime may allocate a *shared* state record for both functions
- HAZARD: The newly allocated executor incidentally retains `this`



# Unintended retention mitigated

---

```
class FileStream {  
  constructor(filename) {  
    let file = new File(filename, "r");  
    this.file = file;  
    openFiles.set(file, makeWeakRef(this, closefile, file));  
  
    // now eagerly load the contents  
    this.loading = file.readAsync().then(data => this.setData(data));  
  } ...  
}
```

- Simple code
- No allocation
- Clean structure pattern