

Weak References for EcmaScript

E. DEAN TRIBBLE – AGORIC SYSTEMS

MARK S. MILLER – GOOGLE, INC.

Background

- The proposed features here are intended for use by library and framework creators, not their clients
- We proposed a version of WeakRefs in 2016
- Bradley Meck made a prototype implementation
- Recent wasm requirement identified some additional issues
- We introduce the notion of a WeakFactory to support manual finalization management and subsystem finalization

Weak References and Finalization

- Fundamental expressiveness
- needed for libraries and frameworks
- to simplify memory management
- and support...
 - Remote references
 - Observers for MVC and data-binding
 - DOM iterators
 - Reactive-style libraries
 - Handles for wasm/external resources

Smalltalk

Java

C#

E

Lua

Haskell

Python

Racket

...

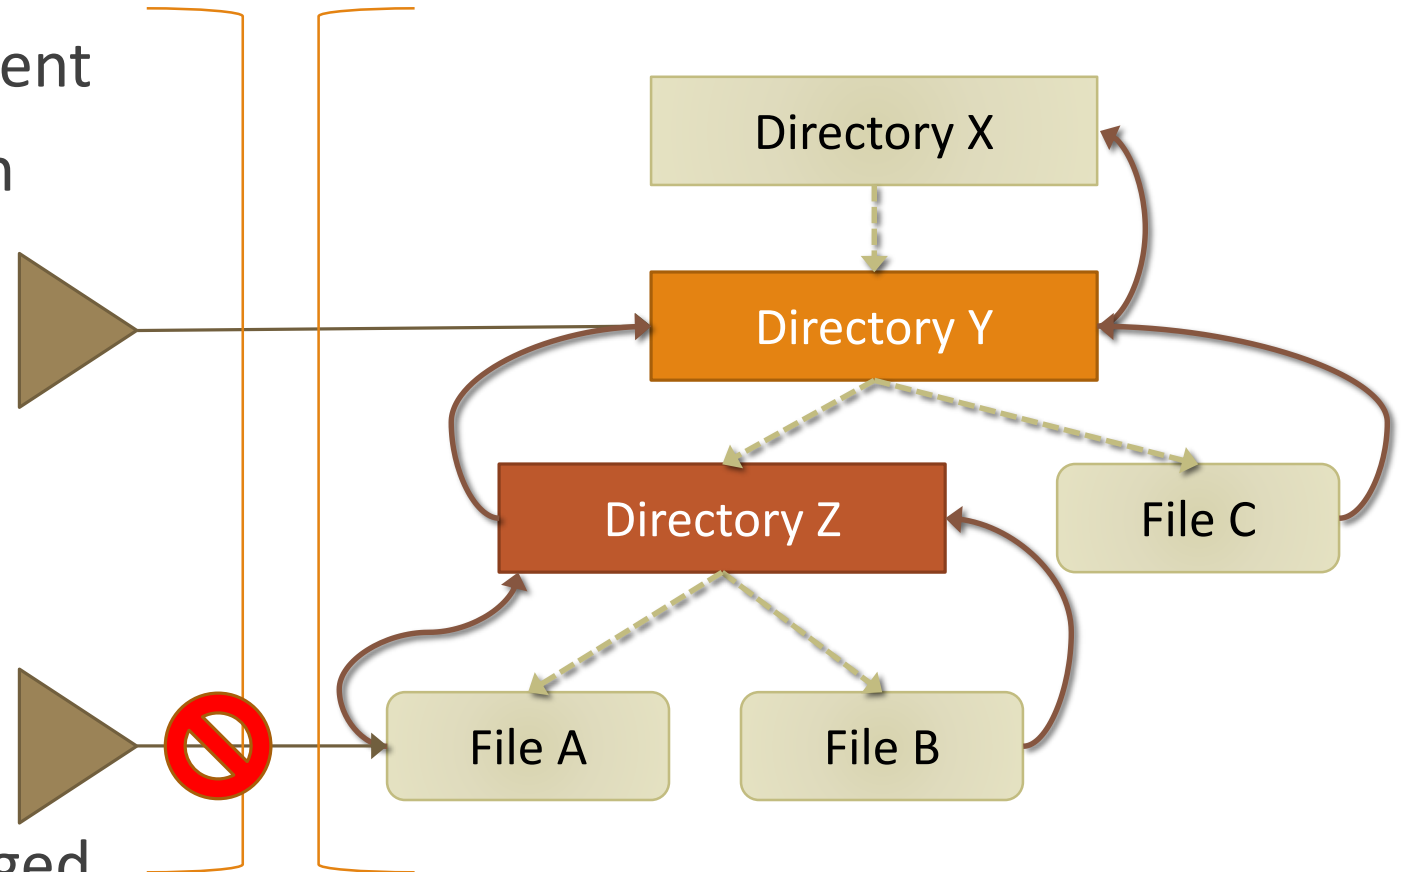
Requirements for Approach

- Multiple, internal and external finalization
 - Remote reference – do something when **I** go away
 - Observers – do something when **you** go away
- Preclude Resurrection
 - Condemned object becomes reachable again
 - But its world is broken
- Avoid Layered collection
 - Data structure is reclaimed one layer at a time

Resurrection Damage

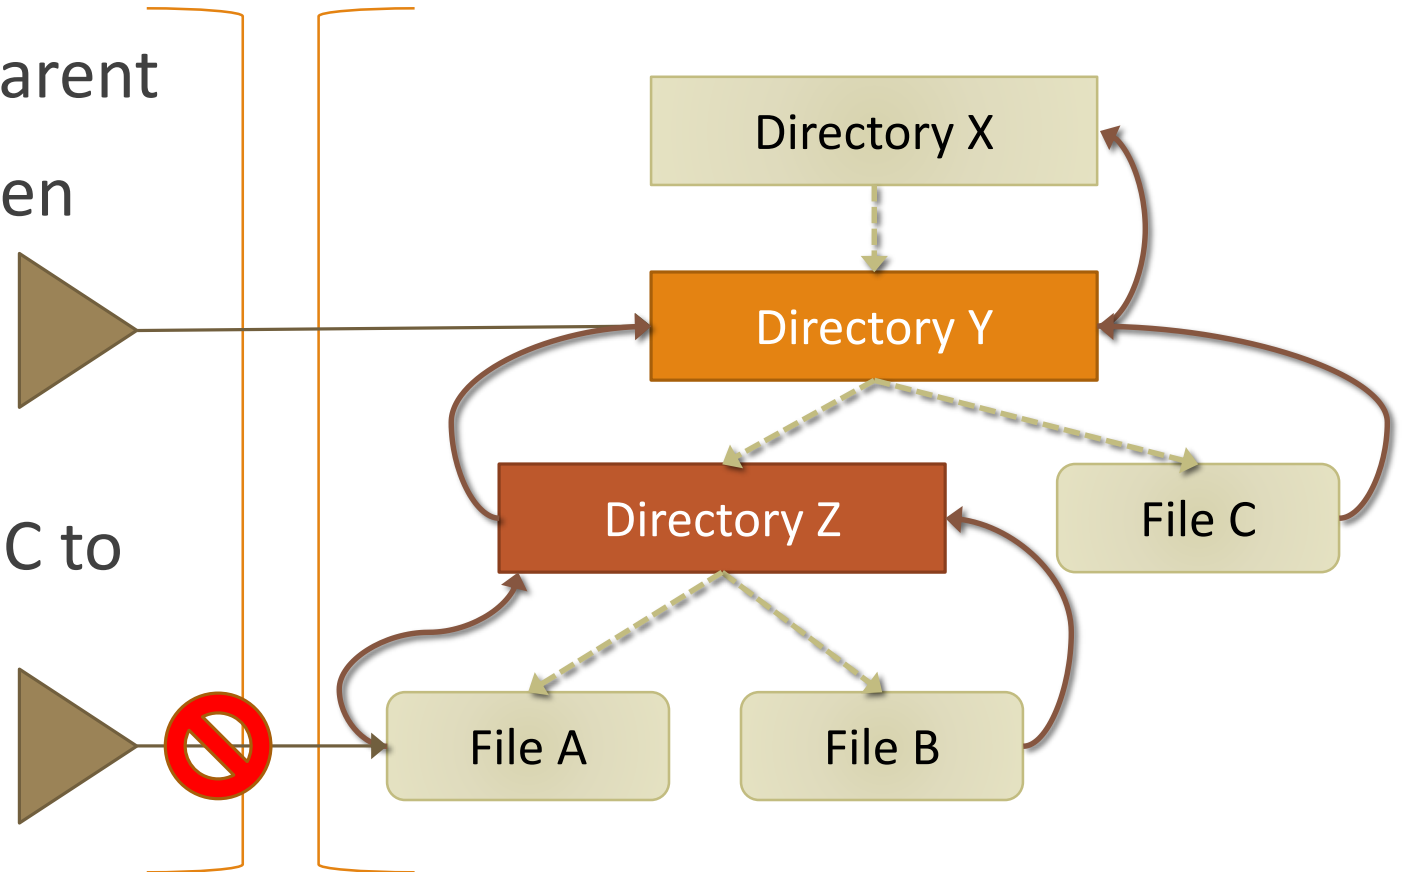
- Files/directories point at parent
- Parent canonicalizes children

- A gets finalized first
- Z is made reachable again
- HAZARD: The Z tree is damaged



Layered Collection

- Files/directories point at parent
 - Parent canonicalizes children
 - Cleanup at each layer
-
- Don't require a separate GC to reclaim each layer

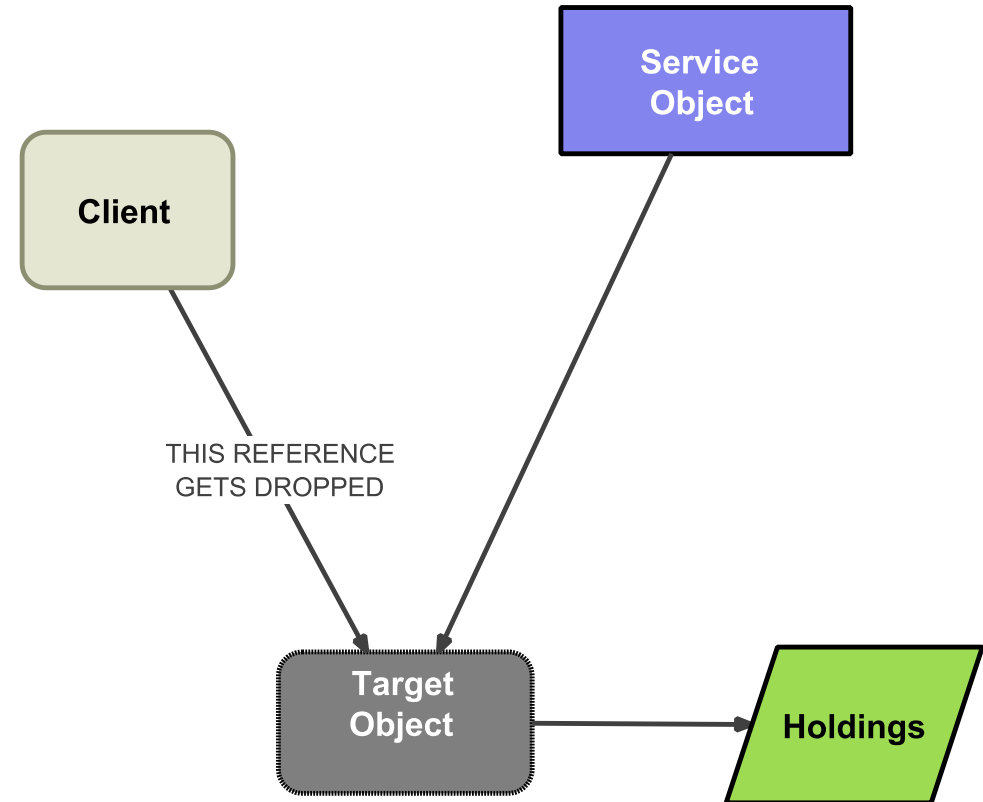


Terminology

- **Reachable** objects – objects that program execution can reach
- **Strongly Reachable** objects – objects that program execution can reach without dereferencing weak references
- **Reclaimed** object – an object that the garbage has noticed is not strongly reachable and has made not reachable
- **Weak reference** – allows access to an object that has not yet been reclaimed, but does not prevent that object from being reclaimed
- **Finalization** – the execution of code to clean up after an object that has been reclaimed
- **Conservative** – some objects that are not strongly reachable may never be reclaimed

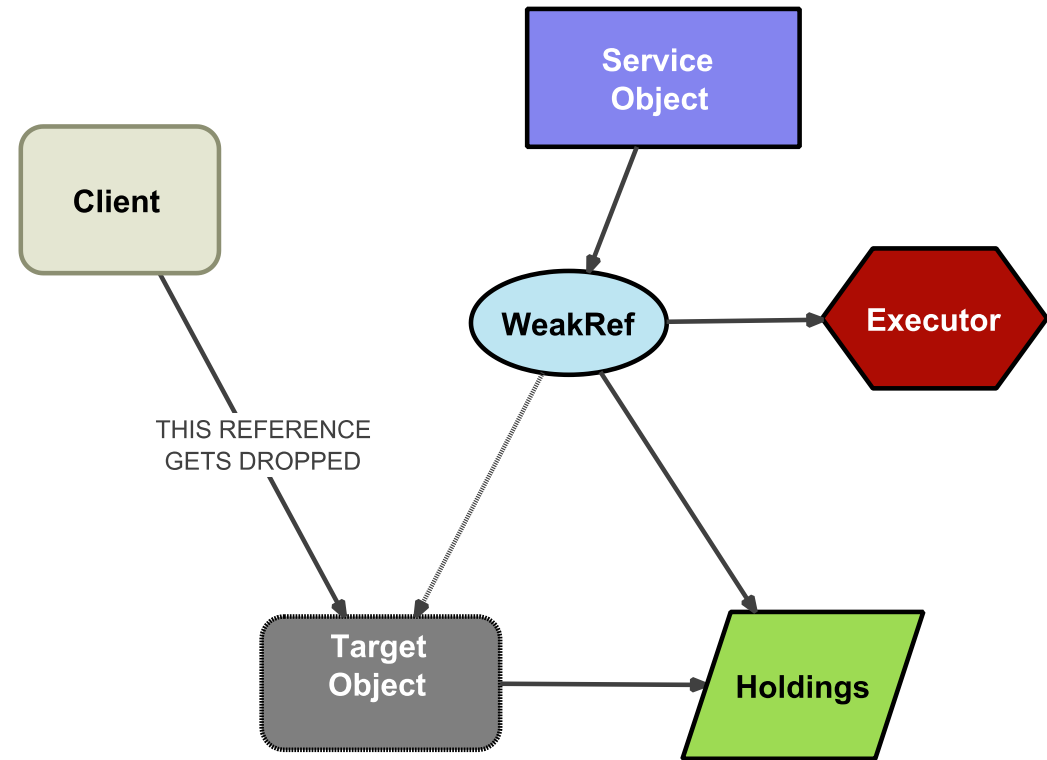
Basic Scenario

- Client gets target from the service
- Target uses holdings to accomplish it's function
- When client drops target, service wants to clean it up
- Service retains Target indefinitely!



Basic Approach

- Insert a WeakRef
- Weak reference shares holdings
- Service's cleanup is scheduled when Target is reclaimed
- The cleanup uses Holdings



Proposed (Partial) API

- **WeakRef**
 - Points weakly at a Target
 - Holds an associated value that gets used during cleanup for a Target
- **WeakFactory**
 - Create `WeakRefs` for a related group of Targets
 - Manage cleanup for `WeakRefs` whose Targets have been reclaimed

class **WeakRef** {

- **deref ()**
 - Return a strong reference to the weakly-held Target object, or undefined if it has been reclaimed
 - The Target will be retained until the end of the turn
- **clear ()**
 - Drop the weak reference to the Target and prevent finalization
- **get holdings ()**
 - Return the holdings associated with the Target

class **WeakFactory** {

- **constructor** (cleanup?: Cleaner)
 - cleanup (weakRefs) => void
 - Invoked in a new job to cleanup after reclaimed targets
 - weakRefs – iterator of weakRefs whose Targets are reclaimed
- **makeRef** (target, holdings?) : WeakRef
 - target – object pointed to weakly; returned by deref ()
 - holdings – target-specific arg used for cleanup after the Target (optional)
- **shutdown** () : void
 - The subsystem is being shutdown; perform no cleanup action for any WeakRef in this group
- **cleanupSome** (cleanupNow : Cleaner) : void

WeakRef States

- Available
 - Target is accessible
- Dirty
 - Target is inaccessible but cleanup is needed
- Clean
 - Target is inaccessible and cleanup is NOT needed

Example: Observer without finalization

- Model points weakly to observers
- Doesn't really cleanup after reclaimed observers

```
class Model {  
    #observers, #weakFactory;  
    constructor() {  
        this.#observers = new Set();  
        this.#weakFactory = new WeakFactory();  
    }  
  
    addObserver(observer) {  
        const wr = this.#weakFactory.makeRef(observer);  
        this.#observers.add(wr);  
    }  
    notify(msg) {  
        this.#observers.forEach(wr =>  
            wr.deref() && wr.deref().changed(msg));  
    }  
}
```


Example: Observer using Finalization

- WeakRefs for reclaimed observers are removed
- No cleanup needed if the whole model is reclaimed

```
class Model {  
  #observers, #weakFactory;  
  constructor() {  
    this.#observers = new Set();  
    const cleanup = iter => {  
      for (const wr of iter) {  
        this.#observers.delete(wr);  
      }  
    };  
    this.#weakFactory = new WeakFactory(cleanup);  
  }  
  addObserver(observer) {  
    const wr = this.#weakFactory.makeRef(observer);  
    this.#observers.add(wr);  
  }  
  notify(msg) {  
    this.#observers.forEach(wr =>  
      wr.deref() && wr.deref().changed(msg));  
  }  
}
```

Example: Observer using Finalization

- Simple code pattern shows a potential hazard



```
class Model {  
  #observers, #weakFactory;  
  constructor() {  
    this.#observers = new Set();  
    const cleanup = iter => {  
      for (const wr of iter) {  
        this.#observers.delete(wr);  
      }  
    };  
    this.#weakFactory = new WeakFactory(cleanup);  
  }  
  addObserver(observer) {  
    const wr = this.#weakFactory.makeRef(observer);  
    this.#observers.add(wr);  
  }  
  notify(msg) {  
    this.#observers.forEach(wr =>  
      wr.deref() && wr.deref().changed(msg));  
  }  
}
```


“Read Consistency”

- The multiple-use hazard

```
weak.deref() && weak.deref().changed(msg);
```

```
if (observers.get(myKey)) {  
    ...do some expensive setup...  
    doOperation(myName, observers.get(myKey));  
}
```

- *A program cannot observe a target getting reclaimed within the execution of a job (turn of the event loop)*
- Trivially precludes the multiple-use hazard
- Minimizes visible non-determinism

Example: Remote Reference

- Client-side reference to an object on the server
- Over a shared connection
- Send “DROP” when the client ref is reclaimed
- No cleanup if the whole connection is reclaimed

```
class RemoteConnection {
  #transport, #remotes, #weakFactory;
  constructor(transport) {
    this.#transport = transport;
    this.#remotes = new Map();
    const cleanup = iter => {
      for (const wr of iter) {
        this.dropRef(wr.holdings);
      }
    };
    this.#weakFactory = new WeakFactory(cleanup);
  }
  makeRemoteRef(remoteId) {
    const remoteRef = ...; // remoteRef construction elided
    const wr = this.#weakFactory.makeRef(remoteRef, remoteId);
    this.#remotes.set(remoteId, wr);
    return remoteRef;
  }
  dropRef(remoteId) {
    this.#transport.send("DROP", remoteId);
    this.#remotes.delete(remoteId);
  }
  ...
}
```

Example: Connecting JS to wasm

- JS wrapper of an unmanaged wasm obj
- Bookkeeping by WasmConnection
- `wasmBridge.delete` when wrapper is reclaimed
- No cleanup needed if wasm instance is reclaimed

```
class WasmConnection {
  #wasmBridge, #wrappers, #cleanup, #weakFactory;
  constructor(wasmBridge) {
    this.#wasmBridge = wasmBridge;
    this.#wrappers = new Map();
    this.#cleanup = iter => {
      for (const wr of iter) {
        this.dropWrapper(wr.holdings);
      }
    };
    this.#weakFactory = new WeakFactory(cleanup);
  }
  makeWrapper(wasmAddr) {
    const wrapper = this.#wasmBridge.makeWrapper(wasmAddr);
    const wr = this.#weakFactory.makeRef(wrapper, wasmAddr);
    this.#wrappers.set(wasmAddr, wr);
    return wrapper;
  }
  dropWrapper(wasmAddr) {
    this.#wasmBridge.delete(wasmAddr);
    this.#wrappers.delete(wasmAddr);
  }
  ...
}
```

Example: wasm without deref ()

- makeRef doesn't work well in a long turn because of read consistency
- RemoteConnection requires canonicalization
- Wasm host binding has not needed that
- So we introduce WeakCell

```
class WasmConnection {
  #wasmBridge, #wrappers, #cleanup, #weakFactory;
  constructor(wasmBridge) {
    this.#wasmBridge = wasmBridge;
    this.#wrappers = new Map();
    this.#cleanup = iter => {
      for (const wr of iter) {
        this.dropWrapper(wr.holdings);
      }
    };
    this.#weakFactory = new WeakFactory(cleanup);
  }
  makeWrapper(wasmAddr) {
    const wrapper = this.#wasmBridge.makeWrapper(wasmAddr);
    const wr = this.#weakFactory.makeRef(wrapper, wasmAddr);
    this.#wrappers.set(wasmAddr, wr);
    return wrapper;
  }
  dropWrapper(wasmAddr) {
    this.#wasmBridge.delete(wasmAddr);
    this.#wrappers.delete(wasmAddr);
  }
  ...
}
```

Introducing **WeakCell**

- **WeakCell**
 - Points weakly at a Target
 - Holds an associated value that gets used during cleanup for a Target
- **WeakRef** extends `WeakCell`
 - Provides ability to dereference Target
 - Creation and dereference preserve the Target til tend of turn
- **WeakFactory**
 - Create `WeakCells` for a related group of Targets
 - Manage cleanup for `WeakCells` whose Targets have been reclaimed

Proposed Complete API

class **WeakCell** {

- **clear()**
 - Drop the weak reference to the Target and prevent finalization
- **get_holdings()**
 - Return the holdings associated with the Target

class **WeakRef** extends WeakCell {

deref()

- Return a strong reference to the weakly-held Target object, or undefined if it has been reclaimed
- The Target will be retained until the end of the turn

class **WeakFactory** {

- **constructor**(cleanup? : Cleaner)
- **makeRef**(target, holdings?) : WeakRef
- **shutdown**() : void
- **makeCell**(target, holdings?) : WeakCell
 - target – object pointed to weakly
 - holdings – target-specific arg used for cleanup after the Target (optional)
- **cleanupSome**(cleanupNow : Cleaner) : void
 - cleanupNow is invoked synchronously to cleanup after some reclaimed targets
- Type **Cleaner** = (weakRefs : Iterator) => void
 - weakRefs – iterator of weakRefs whose Targets are reclaimed

Example: JS ↔ long lived wasm

- **WasmConnection** encapsulates **weakFactory** and all its weakrefs
- **cleanupSome** enables synchronous finalization controlled by the app

```
class WasmConnection {  
  #wasmBridge, #wrappers, #cleanup, #weakFactory;  
  constructor(wasmBridge) { ... }  
  dropWrapper(wasmAddr) { ... }  
  
  makeWrapper(wasmAddr) {  
    const wrapper = this.#wasmBridge.makeWrapper(wasmAddr);  
    const wr = this.#weakFactory.makeCell(wrapper, wasmAddr);  
    this.#wrappers.set(wasmAddr, wr);  
    return wrapper;  
  }  
  
  // Cleanup after reclaimed objects within a turn  
  deleteSome() {  
    this.#weakFactory.cleanupSome(this.#cleanup);  
  }  
}
```


Proposal Characteristics

- Automatic cleanup actions are scheduled in their own jobs
- Multiple, independent WeakRefs could have the same target
 - Internal finalization – remote reference
 - External finalization - observer
- Clients may unregister from finalization (`weakRef.clear()`)

Semantic Details

- WeakCell
 - Points weakly at Target
 - Points strongly at Holdings
 - Points strongly at WeakFactory
- WeakFactory
 - Points strongly at cleanup function
 - Points strongly at Available or Dirty WeakCells in group
- Cleanup iteration
 - Iterator is only productive during the call to cleanup
 - WeakCell is marked clean when it is pulled from the iterator

Non-determinism in GC

- Weak references make GC behavior visible
- Racy-reads/writes present the same problems
 - It can appear sequentially consistent
 - If a program counts on that, it appears correct until it fails in production
- The non-determinism “bandwidth” is limited
- We can further mitigate it

Minimize and contain non-determinism

- Deterministic computation advantages
 - Testability of ES components, frameworks, and applications
 - Reproducibility of results and bug reports
 - Portability across different runtimes and environments
 - Restricts who can read side-channels
- Read consistency
- Weak Reference construction should be closely held
 - `WeakFactory` is in the `System` namespace
- Cross-realm references are strong

Open questions

- Are WeakRefs retained by their WeakFactory?
 - Fallback: Yes, until cleaned
- Is there a query operation for the states of a WeakRef
 - Fallback: No
- `WeakFactory.prototype.shutdown()` could instead use `CancellationToken`

Thank you

WeakValueMap – with Leak

- Keys and Values are weak
- Entry remains after value is GC'd
 - LEAK!

```
class WeakValueMap {  
    constructor() {  
        this.map = new WeakMap();  
    }  
  
    get(key) {  
        let weakRef = this.map[key];  
        return weakRef && weakRef.get();  
    }  
  
    set(key, value) {  
        map.set(key, makeWeakRef(value));  
    }  
}
```

WeakValueMap

- Keys and Values are weak
- Cleanup keys when value is GC'd
- Avoids storage leak

```
class WeakValueMap {  
    constructor() {  
        this.map = new WeakMap();  
        this.executor =  
            keyRef => this.map.delete(keyRef.get());  
    }  
  
    get(key) {  
        let weakRef = this.map[key];  
        return weakRef && weakRef.get();  
    }  
  
    set(key, value) {  
        let keyRef = makeWeakRef(key);  
        let valRef = makeWeakRef(v, this.executor, keyRef);  
        map.set(key, valRef);  
    }  
}
```


More Terminology

- ***Condemned*** object – a (strongly) unreachable object that the garbage collector has noticed it can reclaim

Read Consistency

- A target could be pointed to by multiple WeakRefs

Unintended retention

```
const openfiles = new Map(); // file => weakRef(filestream)

const cleanupFile(file) {
  file.close();
  openFiles.delete(file);
  console.info("Filestream gc before close: ", file.name)
}

class FileStream {
  constructor(filename) {
    this.file = new File(filename, "r");
    openFiles.set(file, makeWeakRef(this, () => closeFile(this.file)));

    // now eagerly load the contents
    this.loading = file.readAsync().then(data => this.setData(data));
  } ...
}
```

- HAZARD: The executor function retains `this`

Unintended retention – runtime black magic

```
class FileStream {  
  constructor(filename) {  
    let file = new File(filename, "r");  
    this.file = file;  
    openFiles.set(file, new WeakRef(this, () => closeFile(file)));  
  
    // now eagerly load the contents  
    this.loading = file.readAsync().then(data => this.setData(data));  
  } ...  
}
```

- The *unrelated* second function closes over `this`
- The runtime may allocate a *shared* state record for both functions
- HAZARD: The newly allocated executor incidentally retains `this`

Unintended retention mitigated

```
class FileStream {  
  constructor(filename) {  
    let file = new File(filename, "r");  
    this.file = file;  
    openFiles.set(file, makeWeakRef(this, closefile, file));  
  
    // now eagerly load the contents  
    this.loading = file.readAsync().then(data => this.setData(data));  
  } ...  
}
```

- Simple code
- No allocation
- Clean structure pattern